Are Your V8 GC Logs Speaking to You?

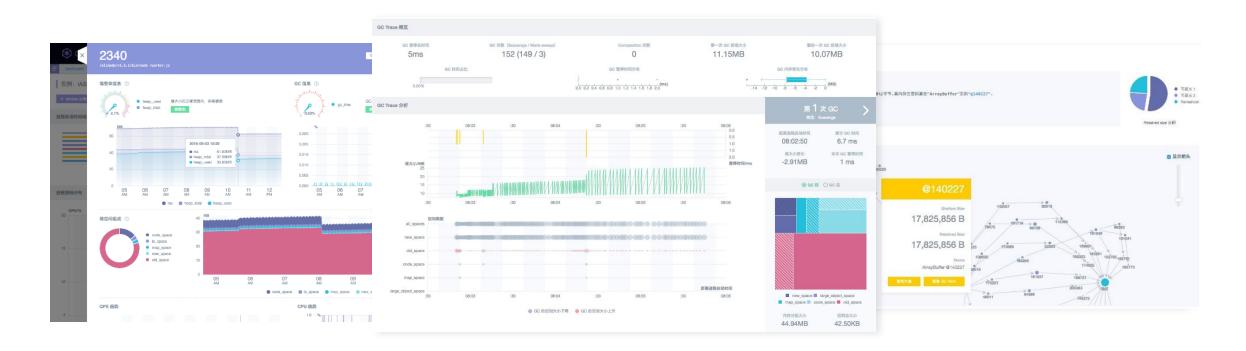
About Me

- @joyeecheung (GitHub, Twitter)
- ❖ Chau Yee Cheung (Cantonese) Qiuyi Zhang /张秋怡 (Mandarin)
- Joyee
- ♦ Intern @

 alinode

 Alibaba Group

 Alibaba



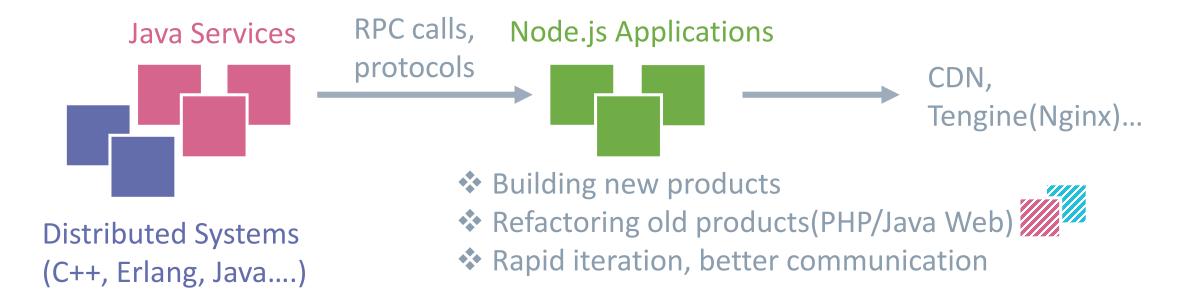
Background

- We provide performance management services to teams both inside and outside the group
- Alibaba is one of the largest companies in China that uses Node.js in production
 - Mostly driven by frontend devs



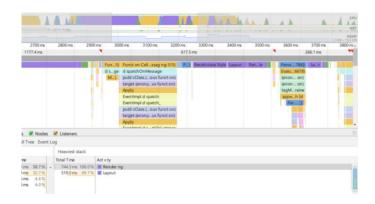




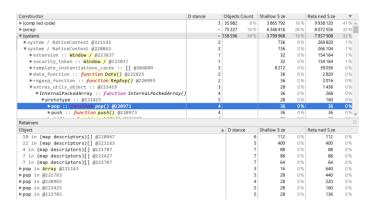


Background

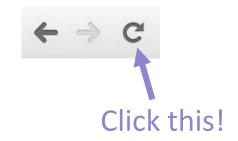
- Java architects argue that Node.js is not mature enough
 - Lack of tooling for monitoring and profiling
- Java programmers know their VM and their GC. As for JS programmers, well...
 - Most come from a frontend dev background



CPU profile
The pursuit of 60fps



HeapSnapshopts / HeapTimelines you don't want to blow up your users' RAM

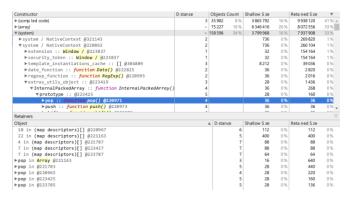


GC logs
Servers are long-running,
clients are not

❖ V8 garbage collection logs are the least documented tools, not even available in devtools

What are V8 Garbage Collection logs good for?

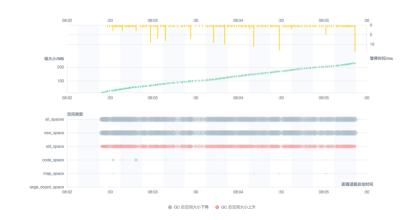
Combined with heap profiles, they can help you catch the culprit for memory leaks



Heapdump can tell you what is leaking, but

they are sometimes too specific





Help you verify that your fixes do fix your leaks

- What is Garbage Collection(GC)?
 - Collect and free your objects(garbage) when you don't need them anymore. That's why you don't need to manage the memory yourself in languages like JavaScript, Java, Ruby...
 - For ECMAScript/JavaScript, there are no specifications about it, everything is up to the implementation
- This talk is based on V8 4.5.103.35 (the one used by Node v4.x)
 - Node v6.x will use V8 5.x, which introduces a few improvements(mostly parallel stuff) and changes the GC log format a bit, but most of this talk still apply

- Generational hypothesis
 - Most objects die young, others tend to live forever
- ❖ V8 GC Strategy
 - Generational: Young and Old
 - Accurate: Given a random word(always at least 4-byte aligned) in the memory, V8 can tell you if it's a pointer or some data by looking at the last bit reserved as a tag
 - So V8 doesn't have to scan the memory to be sure about this, which makes updating the pointers during GC fairly sufficient.
- ❖ GC is triggered by allocations(new or var most of the time), when the memory allocated for a space is not enough
- ❖ In V8, Objects live on the heap, and the heap is divided into spaces

Code Space


Executable code compiled by V8

Map Space

Metadata of hidden classes, pointed by objects

Large Object Space

Objects too large to fit into any other spaces.



Old Space

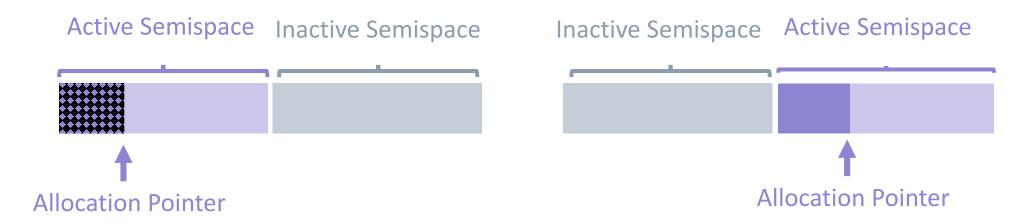
Objects survive more than two GCs in the new space

New Space

Divided in half, only one is active at a time.

Most objects start and end their lives here.

- New space: Scavenger
 - Implementation of Cheney's algorithm



- Old space
 - Mark-sweep/mark-compact
- Oddly similar to the early HotSpot JVM(designed by the same person)
- Further reading
 - http://alinode.aliyun.com/blog/14 (Chinese)
 - http://jayconrod.com/posts/55/a-tour-of-v8-garbage-collection (English)

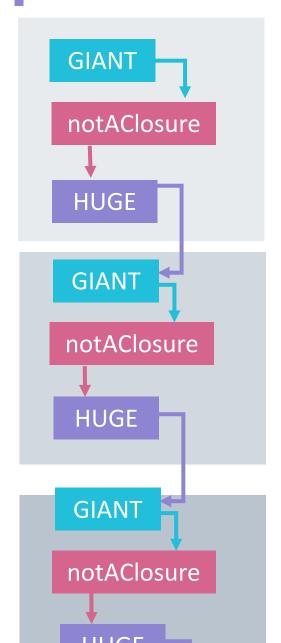
Getting started with V8 GC logs

- Plenty of options(node --v8-options| grep gc)
- We'll be focusing on
 - --trace_gc
 - --trace_gc_nvp
- Obviously you can't use these in production, even though it would be difficult to reproduce your problems offline
 - v8.setFlagsFromString('--trace_gc')
 - v8.setFlagsFromString('--notrace_gc')
 - Or use alinode to get them by clicking a button on our SAAS platform(wink wink)
- In case you do want to read the code
 - deps/v8/src/heap, Start With gc-tracer.cc

Demo Time

- Closures are created when declared, not when executed
- In V8, once there is any closure in the context, the context will be attached to every function, even for those who don't reference the context at all

```
const express = require('express');
const app = express();
var GIANT;
function leak() {
  var HUGE = GIANT;
  function unusedClosure() {
    HUGE.slice(1);
  GIANT = {
    willBeLeaked: new Array(1e5).join('.'),
    notAClosure: function notAClosure() {
      return 1;
app.get('/', function handler(req, res) {
  leak();
  console.log(new Date());
  res.send('Hello World!');
app.listen(3000, function startApp() {
  console.log('Example app listening on port 3000!');
});
```



Created in this call

Created in last call, but can't be GC'ed because they are referenced by the new ones

```
const express = require('express');
const app = express();
var GIANT;
                        This shouldn't have happened!
function leak()
  var HUGE = GIANT;
  function unusedClosure() {
                                       It happens
    HUGE.slice(1);
                                       because of
                                       this
  GIANT = {
    willBeLeaked: newVArray(1e5).join('.'),
    notAClosure: function notAClosure() {
      return 1;
app.get('/', function handler(req, res) {
  leak();
  console.log(new Date());
  res.send('Hello World!');
});
app.listen(3000, function startApp() {
  console.log('Example app listening on port 3000!');
});
```

Constructor	Distance	Objects C	Shallow Size	F	Retained Si	ize ▼
▼(string)	3	723 27%	031696 9	8% 0	31696	98%
"	6172		100 024	0% 1	100 024	0%
"	6175		100 024	0% 1	100 024	0 %
"	3988		100 024	0% 1	100 024	0 %
"	4372		100 024	0% 1	100 024	0 %
"	4375		100 024	0% 1	100 024	0 %
"	3985		100 024	0% 1	100 024	0 %
"	3619		100 024	0% 1	100 024	0 %
"	3622		100 024	0% 1	100 024	0 %
"	2806		100 024	0% 1	100 024	0 %
"	2809		100 024	0% 1	100 024	0 %
Retainers						=
Object		Distance ▲	Shallow Size	F	Retained Si	ize
▼willBeLeaked in @13957		6171	40	0 %	40	0 %
▼HUGE in system / Context @100513		6170	56	0 %	56	0 %
▼context in function() @100511		6169	72	0 %	72	0 %
▼notAClosure in @100507		6168	40	0 %	40	0 %
▼HUGE in system / Context @100505		6167	56	0 %	56	0 %
▼context in function() @100503		6166	72	0 %	72	0 %
▼notAClosure in @100499		6165	40	0 %	40	0 %
▼HUGE in system / Context @100497		6164	56	0 %	56	0 %
▼context in function() @100495		6163	72	0 %	72	0 %
▼notAClosure in @100491		6162	40	0 %	40	0 %
▼HUGE in system / Context @100489		6161	56	0 %	56	0 %
▼context in function() @100487		6160	72	0 %	72	0 %
▼notAClosure in @100483		6159	40	0 %	40	0 %
▼HUGE in system / Context @100481		6158	56	0 %	56	0 %
▼context in function() @100479		6157	72	0 %	72	0 %
▼notAClosure in @100475		6156	40	0 %	40	0 %

08:02

:30

08:03

Spent on --trace gc (print one trace line following each garbage collection) external Size of memory allocated from OS memory The time since you start this VM 611604 ms: Scavenge 1153.6 [1457.9] -> 1152.9 [1457.9] MB, 0.6 / 0 ms [allocation failure]. 19642:0x102004a00 611774 ms: Scavenge 1153.7 (1457.9) -> 1152.9 (1457.9) MB, 0.6 / 0 ms [allocation failure]. [19642:0x102004a00] [19642:0x102004a00] 611775 ms: Scavenge 1153.7 (1457.9) -> 1153.0 (1457.9) MB, 0.6 / 0 ms [allocation failure]. 611880 ms: Scavenge 1153.1 (1457.9) -> 1153.1 (1457.9) MB, 0.5 / 0 ms [allocation failure]. [19642:0x102004a00] 611881 ms: Scavenge 1153.1 (1457.9) -> 1153.1 (1457.9) MB, 0.5 / 0 ms [allocation failure] [19642:0x102004a00] [19642:0x102004a00] 611916 ms: Mark-sweep 1153.1 (1457.9) -> 1152.4 (1449.9) MB. 35.0 / 0 ms [last resort gc]. 611951 ms: Mark-sweep 1152.4 (1449.9) -> 1152.3 (1442.9) MB 34.8 0 ms [last resort gc]. [19642:0x102004a00] Type of GC Size of all objects Pause Because the GC is triggered when 08:02 :30 08:03 08:04 08:05 there's not enough memory 10 堆大小/MB 暂停时间/ms 100

08:04

:30

08:05

:30

:30

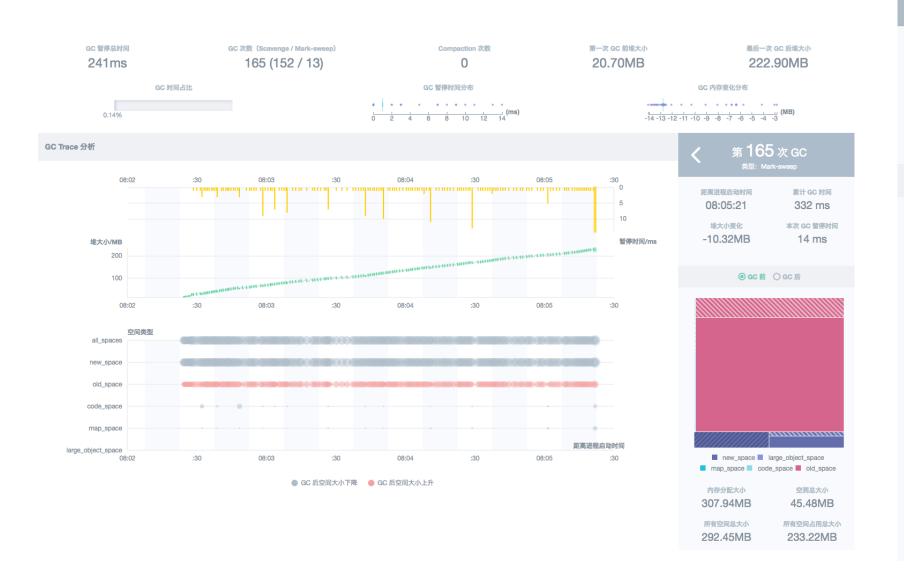
 --trace_gc_nvp (print one detailed trace line in name=value format after each garbage collection)

Name-value

Type of GC, ms=Mark-sweep/Mark-compact, s=Scavenger

```
[12543:0x102004a00] [I:0x102004a00] 152163 ms: pause=3.2 mutator=729.4 gc=ms external=0.0 mark=0.3 sweep=2.27
sweepns=1.53 sweepos=0.15 sweepcode=0.18 sweepcell=0.00 sweepmap=0.05 evacuate=0.0 new new=0.0 root new=0.0
old_new=0.0 compaction_ptrs=0.0 intracompaction_ptrs=0.0 misc_compaction=0.3 weak_closure=0.0 inc_weak_closure=0.0
weakcollection process=0.0 weakcollection clear=0.0 weakcollection abort=0.0 total size before=37492960
total size after=25744424 holes size before=2600536 holes size after=3375144 allocated=11756832 promoted=1407928
semi space copied=1119104 nodes died in new=12 nodes copied in new=1 nodes promoted=0 promotion ratio=10.6%
average_survival_ratio=17.5% promotion_rate=97.4% semi_space_copy_rate=8.4% new_space_allocation_throughput=14433
context disposal rate=0.0 steps count=51 steps took=7.2 longest step=0.9 incremental marking throughput=2560955
[12543:0x102004a00] Memory allocator, used: 65476 KB, available: 1433660 KB
[12543:0x102004a00] New space,
                                 used:
                                       1105 KB, available: 15018 KB, committed: 32248 KB
[12543:0x102004a00] Old space,
                                 used: 22174 KB, available:
                                                              2622 KB, committed: 25049 KB
                                                                                             Allocatable memory
[12543:0x102004a00] Code space,
                                 used: 1617 KB, available:
                                                              622 KB, committed:
                                                                                    2266 KB
                                                                                             and page headers
[12543:0x102004a00] Map space, used: 243 KB, available: 810 KB, committed:
                                                                                    1070 KB
                                                 0 KB, available: 1432619 KB, committed:
[12543:0x102004a00] Large object space, used:
                                                                                             0 KB
[12543:0x102004a00] All spaces,
                                      used: 25141 KB, available: 1451694 KB, committed: 60634 KB
[12543:0x102004a00] External memory reported: 16 KB Memory left in the allocator
                                                           committed = (available + used)(*2 if it's new space)
[12543:0x102004a00] Total time spent in GC : 32.5 ms
```





第 **1** 次 GC

类型: Scavenge

距离进程启动时间

累计 GC 时间

08:02:24

20.1 ms


堆大小变化

本次 GC 暂停时间

-6.79MB

0 ms

● GC前 ○ GC后



```
FATAL ERROR: CALL_AND_RETRY_LAST Allocation failed - process out of memory Abort trap: 6
```

deps/v8/src/api.cc: ResourceConstraints::ConfigureDefaults

```
-const uint64_t low_limit = 512ul * i::MB;
-const uint64_t medium_limit = 768ul * i::MB;
-const uint64_t high_limit = 1ul * i::GB;
```

deps/v8/src/heap/heap.h

```
// The old space size has to be a multiple of Page::kPageSize.
// Sizes are in MB.
static const int kMaxOldSpaceSizeLowMemoryDevice = 128 * kPointerMultiplier;
static const int kMaxOldSpaceSizeMediumMemoryDevice =
....256 * kPointerMultiplier;
static const int kMaxOldSpaceSizeHighMemoryDevice = 512 * kPointerMultiplier;
static const int kMaxOldSpaceSizeHugeMemoryDevice = 700 * kPointerMultiplier;
static const int kPointerMultiplier = i::kPointerSize / 4;
const int kPointerSize = sizeof(void*); // NOLINT
```



Solution #1: nullify your reference

Cut off the reference to the old GIANT when you don't need it anymore

```
var GIANT;
function leak() {
  var HUGE = GIANT;
  function unusedClosure() {
   HUGE.slice(1);
  GIANT = {
   willBeLeaked: new Array(1e5).join('.'),
    notAClosure: function notAClosure() {
      return 1;
 HUGE = null; /* not used anymore! */
```

Solution #1: nullify your reference



Solution #2: passing parameters

Avoid closure like a plague

```
var GIANT;
function leak() {
  var HUGE = GIANT;
  function notAClosureAnymore(HUGE_ARG) {
    HUGE_ARG.slice(1);
  GIANT = {
    willBeLeaked: new Array(1e5).join('.'),
    notAClosure: function notAClosure() {
      return 1;
```

Solution #2: passing parameters



Questions?



https://alinode.aliyun.com qiuyi.zqy@alibaba-inc.com @joyeecheung (GitHub, Twitter)